

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import MultinomialNB
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
```

```
In [3]: data = {
    "label": ["spam", "ham", "spam", "ham", "spam", "ham", "spam", "ham"],
    "message": [
        "Win a free iPhone now",
        "Hey, are we meeting today?",
        "Congratulations, you won a prize",
        "Can we talk later?",
        "Claim your free reward now",
        "See you at dinner",
        "You have won lottery money",
        "Let's go to the movie"
    ]
}

df = pd.DataFrame(data)
df
```

	label	message
0	spam	Win a free iPhone now
1	ham	Hey, are we meeting today?
2	spam	Congratulations, you won a prize
3	ham	Can we talk later?
4	spam	Claim your free reward now
5	ham	See you at dinner
6	spam	You have won lottery money
7	ham	Let's go to the movie

```
In [4]: df['label'] = df['label'].map({'spam': 1, 'ham': 0})
df
```

Out[4]:

	label	message
0	1	Win a free iPhone now
1	0	Hey, are we meeting today?
2	1	Congratulations, you won a prize
3	0	Can we talk later?
4	1	Claim your free reward now
5	0	See you at dinner
6	1	You have won lottery money
7	0	Let's go to the movie

In [5]:

```
vectorizer = TfidfVectorizer(stop_words='english')
X = vectorizer.fit_transform(df['message'])
y = df['label']
```

In [6]:

```
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.25, random_state=42
)
```

In [7]:

```
model = MultinomialNB()
model.fit(X_train, y_train)
```

Out[7]:

▼ MultinomialNB ⓘ ⓘ

► Parameters

In [8]:

```
y_pred = model.predict(X_test)
y_pred
```

Out[8]:

```
array([1, 1])
```

In [9]:

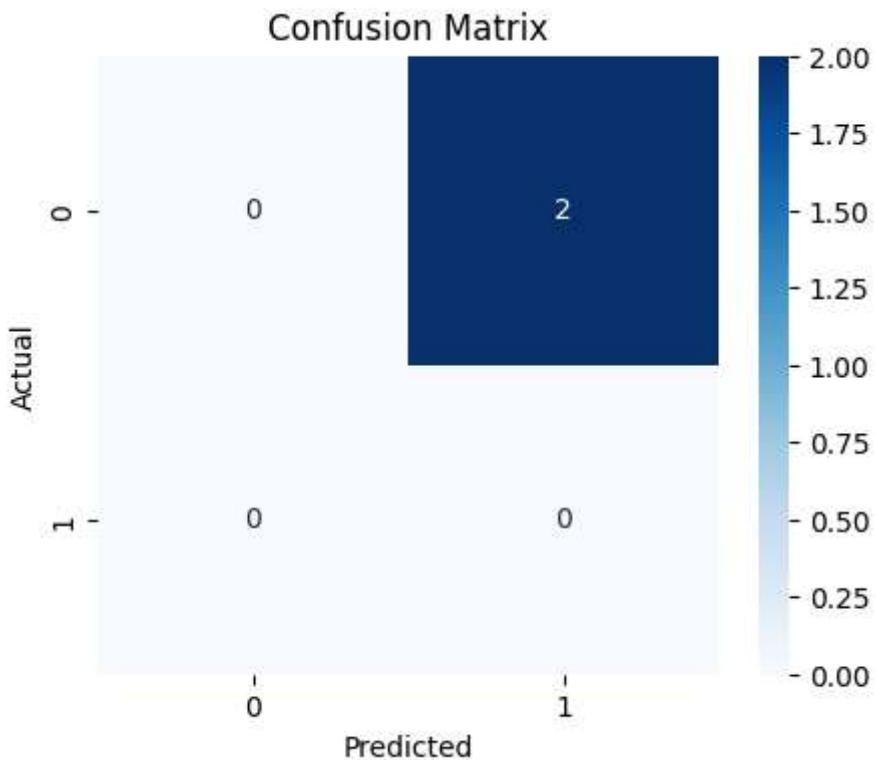
```
accuracy = accuracy_score(y_test, y_pred)
print("Model Accuracy:", accuracy)
```

Model Accuracy: 0.0

In [10]:

```
cm = confusion_matrix(y_test, y_pred)

plt.figure(figsize=(5,4))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix")
plt.show()
```



```
In [11]: print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.00	0.00	0.00	2.0
1	0.00	0.00	0.00	0.0
accuracy			0.00	2.0
macro avg	0.00	0.00	0.00	2.0
weighted avg	0.00	0.00	0.00	2.0

```
C:\Users\Atharav\AppData\Local\Programs\Python\Python313\Lib\site-packages\sklearn\metrics\_classification.py:1833: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", result.shape[0])
C:\Users\Atharav\AppData\Local\Programs\Python\Python313\Lib\site-packages\sklearn\metrics\_classification.py:1833: UndefinedMetricWarning: Recall is ill-defined and being set to 0.0 in labels with no true samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", result.shape[0])
C:\Users\Atharav\AppData\Local\Programs\Python\Python313\Lib\site-packages\sklearn\metrics\_classification.py:1833: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", result.shape[0])
C:\Users\Atharav\AppData\Local\Programs\Python\Python313\Lib\site-packages\sklearn\metrics\_classification.py:1833: UndefinedMetricWarning: Recall is ill-defined and being set to 0.0 in labels with no true samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", result.shape[0])
C:\Users\Atharav\AppData\Local\Programs\Python\Python313\Lib\site-packages\sklearn\metrics\_classification.py:1833: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", result.shape[0])
C:\Users\Atharav\AppData\Local\Programs\Python\Python313\Lib\site-packages\sklearn\metrics\_classification.py:1833: UndefinedMetricWarning: Recall is ill-defined and being set to 0.0 in labels with no true samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", result.shape[0])
```

```
In [13]: #testing
new_email = ["Congratulations! You won free money"]
new_email_vector = vectorizer.transform(new_email)

prediction = model.predict(new_email_vector)

if prediction[0] == 1:
    print("⚠️ Spam Email Detected")
else:
    print("✅ Not Spam Email")
```

⚠️ Spam Email Detected